

# rnn: a Recurrent Neural Network in R<sup>\*</sup>

Bastiaan Quast<sup>†</sup>

19th April 2016

## Abstract

The `rnn` package implements a Recurrent Neural Network (RNN). RNN algorithms have the ability to train neural networks to deal with greater levels of complexity . This package is purposely designed to demonstrate the self learning ability using the classic example of binary summation on a bit-by-bit (right to left) basis, which requires the model to develop the understanding that if a 1 and a 1 are added, the outcome is 0, but in the next iteration, it has to that it was carrying a 1 from the previous iteration.

---

<sup>\*</sup><https://cran.r-project.org/package=rnn> | <https://github.com/bquast/rnn>

<sup>†</sup><http://qua.st> | [bastiaan.quast@graduateinstitute.ch](mailto:bastiaan.quast@graduateinstitute.ch) | [bquast@gmail.com](mailto:bquast@gmail.com)

# 1 Introduction

This package implements a Recurrent Neural Network which is trained to sum 8-bit binary numbers, teaching itself the complex task of carrying a 1 over to the next iteration if the sum of a column takes two bits of space.

to convert numbers in range of 0-127 to binary representation.

Of course, numbers  $< 128$  can be represent in a 7-bit binary form, but since we are adding two numbers in the range 0-127, the total can reach and achieve 128, which requires 8 bits, it cannot be more than 254, the limit of 8 bit binary representation is 255, thereby preventing overflows.

At this point it is useful to clarrify the nomenclature in this article. I use the term RNN (capitalised) for the general concept of a Recurrent Neural Network and I use `rnn` (in miniscules and using a monospace font) to refer to the R package.

Table 1: `rnn` Package

```
# load the package
library(rnn)

# list functions
ls('package:rnn')

## [1] "bin2int"                "int2bin"
## [3] "predictr"               "sigmoid"
## [5] "sigmoid_output_to_derivative" "trainr"
```

As is listed above, the package contains the following functions:

- `bin2int()`: conversion of a matrix of numbers in binary representation to decimal representation;
- `int2bin()`: conversion of a vector numbers in decimal representation to binary representation;
- `predictr()`: predicts response variable based on a `trainr()` model and input data;
- `sigmoid()`: converts any number to a probability between 0 and 1;
- `sigmoid_output_to_derivative()`: takes output of `sigmoid()` and returns the point derivative of that output;
- `trainr()`: primary function, trains a model based on training data and hyperparameters.

In addition to these functions there are also two internal functions `i2b()` and `b2i()`, these functions are used by `int2bin()` and `bin2int()` internally to change a single number from decimal to binary or visa versa.

## 2 Data

The main `trainr()` function takes three integer vectors as inputs: `Y`, `X1`, and `X2`. The vectors `X1` and `X2` are independent variables, the `Y` vector is the sum of `X1` and `X2` and acts as the response variable (for more info see `help('trainr')`).

Training data can be generated using `base` package's `sample()` function. For reproducibility, we also set the seed value of the psuedo-random number generator that `R` uses internally to 1. After generating `X1` and `X2`, I add the two pairwise and store the result in `Y`. Finally, I convert both the input variables and the response variable to binary representation using the `int2bin()` included with the package.

Table 2: Training Data

```
# use the same random numbers
set.seed(1)

# create training inputs
X1 = sample(0:127, 7000, replace=TRUE)
X2 = sample(0:127, 7000, replace=TRUE)

# create training output
Y <- X1 + X2
```

Internally the `int2bin()` function converts these characters into binary format using the `intToBits()` function, the `bin2int()` function converts it back into decimal format for printing using the `packBits()` function, both functions are included in the `base` package.

We can for instance take the first value of `X1` and convert it to a binary representation, whereby the `binary_dim` argument to the `trainr()` function determines the length of the binary representation, throughout this paper we will use 8 bit representations (which limits numbers to the range 0-255), but the theoretical limit is 32 bits.

Table 3: Binary Representation

```
int2bin( X1[1] )

##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
## [1,]    0    0    1    0    0    0    0    1
```

Lets check look at the first sum in decimal representation.

Table 4: Decimal Summation

```
X1[1]

## [1] 33

X2[1]

## [1] 89

X1[1] + X2[1]

## [1] 122

Y[1]

## [1] 122
```

and now in binary representation.

Table 5: Binary Summation

```
int2bin( X1[1] )
int2bin( X2[1] )
print('-----')
int2bin( Y[1] )

##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
## [1,]    0    0    1    0    0    0    0    1
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
## [1,]    0    1    0    1    1    0    0    1
## [1] "-----"
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
## [1,]    0    1    1    1    1    0    1    0
```

As can be seen from the above output, the first values of **X1** and **X2**, 33 and 89 respectively, are both in the range 0-127, which can be represented with only 7 bits. Yet the sum of the two - 122 - is almost outside of the range 0-127, which is why an 8th bit is required (i.e. the 8th value from right to left in the bottom row is 1). If we sampled numbers great than 127 for **X1** and **X2** then the sum of the two could be greater than 255, which requires a ninth bit (or `length=9`).

We can now convert the entire vectors to binary matrices.

```
# convert to binaries of 8 bit (default)
X1 <- int2bin(X1)
X2 <- int2bin(X2)
Y  <- int2bin(Y)
```

### 3 Methodology

The workhorse of the **rnn** package is the **trainr()** function.

For example, if we add the binary numbers 0 0 1 (decimal system: 1) and 1 0 1 (decimal system: 5), we start by adding the right column, 1 and 1 make 1 0 (similar to when 5 and 5 make 1 0 in the decimal system) , the 0 is stored in the right column, the 1 is carried over to the middle column and added with the two existing bits 0 and 0, to form 1, which is stored in the middle column. This time nothing is carried over and the left column sums 0 and 1 to make 1, which gives the outcome 1 1 0 (decimal system: 6).

If we go back to the output of the **int2bin()** function for **X1**, **X2**, and **Y**, we see that in the 4th column (from right to left), a 0 and a 0 are added, resulting in an output of 1. This is because in the previous iteration 3rd column (from right to left) a 1 and a 1 are added, which becomes 1 0, so the 0 goes into column 3 and the 1 is carried over to column 4. Since the summation is done bit by bit (or column by column), the neural network need to remember from the 3rd iteration until the 4th iteration that it is carrying a 1 over. It is this remembering that a feed-forward neural network cannot teach itself.

The **trainr()** and **predictr()** functions internally make use of the **sigmoid()** function, which is a very simple implementation of a sigmoid which takes the range (-Infinity, Infinity) and maps it to the range (0, 1).

Table 6: Sigmoid Source Code

```
# print source code of the sigmoid function
sigmoid

## function(x) {
##   output = 1 / (1+exp(-x))
##   return(output)
## }
## <environment: namespace:rnn>
```

For instance:

Table 7: Sigmoid Examples

```
sigmoid(-137)
sigmoid(5.3)

## [1] 3.174359e-60
## [1] 0.9950332
```


The shape of the sigmoid function is roughly shown below.

Figure 1: Sigmoid Shape

```
library(ggplot2) # load plotting package

# sequence of -10 through 10
x = seq(-10, 10)

# plot sigmoid shape
qplot(x = x, y = sigmoid(x), geom='line')
```



Additionally the `trainr()` and `predictr()` functions use the `sigmoid_output_to_derivative()` function.

Table 8: Sigmoid Derivative Source Code

```
# print source code of the sigmoid_output_to_derivate function
sigmoid_output_to_derivative

## function(output) {
##   return( output*(1-output) )
## <environment: namespace:rnn>
```

As the purpose of the package is to illustrate the working of a Recurrent Neural Network, the `trainr()` function is quite verbose (this can be controlled using the `print` argument).

Table 9: `trainr()` Output

```
## Error in 1:dim(Y)[1]: argument of length 0
```

The `trainr()` function will run until it has evaluated all rows in the matrices that it is fed. Since the training of the network, particularly the carrying part, takes many iterations to learn (the exact number of iterations varies but depends on the hyperparameters, more on this in the next section), it is therefore advisable to sample several thousand values (I use 7000).

The text printed here is of the 8 steps of the summation of the 1000th value of `X1` and `X2`, or iteration 7993-8000.

Each iteration is printed individually, with the two input bits, the prediction for the response value and the actual response value.

After each iteration the difference between the predicted value and the actual value is fed back into the neural network using a method called back-propagation (an application the chain rule of differential calculus).

At the end of the 8 iterations that it here takes to add two values of `X1` and `X2`, the results are printed in a more human legible form. It should be clear from the results that after 1000 numbers, which 8 iterations each, the model is still performing very poorly.

However, progress can be observed:



Table 10: trainr() Output

```
# use the same random numbers
set.seed(1)

# train the network
m1 <- trainr(Y,
             X1,
             X2,
             binary_dim = 8,
             alpha       = 0.1,
             input_dim   = 2,
             hidden_dim  = 10,
             output_dim  = 1,
             print        = 'minimal')

## [1] "Summation number: 1000"
## [1] "Error: 3.96355932085217"
## [1] "X1[ 1000 ]: 0 0 1 0 0 0 0 1 ( 33 )"
## [1] "X2[ 1000 ]: 0 1 1 0 0 1 0 0 + ( 100 )"
## [1] "-----"
## [1] "Y[ 1000 ]: 1 0 0 0 0 1 0 1 ( 133 )"
## [1] "predict Y~: 1 1 0 0 1 1 0 1 ( 205 )"
## [1] "-----"
## [1] "Summation number: 2000"
## [1] "Error: 3.41811756435133"
## [1] "X1[ 2000 ]: 0 0 1 1 0 0 0 1 ( 49 )"
## [1] "X2[ 2000 ]: 0 0 1 1 0 0 1 1 + ( 51 )"
## [1] "-----"
## [1] "Y[ 2000 ]: 0 1 1 0 0 1 0 0 ( 100 )"
## [1] "predict Y~: 0 0 0 0 0 0 1 0 ( 2 )"
## [1] "-----"
## [1] "Summation number: 3000"
## [1] "Error: 3.20534709314338"
## [1] "X1[ 3000 ]: 0 1 0 0 0 1 0 0 ( 68 )"
## [1] "X2[ 3000 ]: 0 1 0 1 0 0 1 1 + ( 83 )"
## [1] "-----"
## [1] "Y[ 3000 ]: 1 0 0 1 0 1 1 1 ( 151 )"
## [1] "predict Y~: 1 0 0 1 1 1 1 1 ( 159 )"
## [1] "-----"
## [1] "Summation number: 4000"
## [1] "Error: 3.44514161374802"
## [1] "X1[ 4000 ]: 0 1 0 1 0 1 0 1 ( 85 )"
## [1] "X2[ 4000 ]: 0 0 1 1 0 1 1 1 + ( 65 )"
## [1] "-----"
## [1] "Y[ 4000 ]: 1 0 0 0 1 1 0 0 ( 140 )"
## [1] "predict Y~: 1 1 1 0 1 0 0 0 ( 232 )"
## [1] "-----"
## [1] "Summation number: 5000"
## [1] "Error: 2.39543018899181"
## [1] "X1[ 5000 ]: 0 0 1 0 0 1 1 0 ( 38 )"
## [1] "X2[ 5000 ]: 0 0 1 0 1 0 1 1 + ( 43 )"
## [1] "-----"
## [1] "Y[ 5000 ]: 0 1 0 1 0 0 0 1 ( 81 )"
## [1] "predict Y~: 0 1 0 1 1 0 0 1 ( 89 )"
## [1] "-----"
## [1] "Summation number: 6000"
## [1] "Error: 1.36041534642064"
## [1] "X1[ 6000 ]: 0 0 1 0 0 1 0 1 ( 37 )"
## [1] "X2[ 6000 ]: 0 1 1 0 1 1 0 1 + ( 109 )"
## [1] "-----"
## [1] "Y[ 6000 ]: 1 0 0 1 0 0 1 0 ( 146 )"
## [1] "predict Y~: 1 0 0 1 0 0 1 0 ( 146 )"
## [1] "-----"
## [1] "Summation number: 7000"
## [1] "Error: 1.02655734929111"
## [1] "X1[ 7000 ]: 0 0 1 0 1 1 0 1 ( 45 )"
## [1] "X2[ 7000 ]: 0 1 0 1 1 1 0 0 + ( 92 )"
## [1] "-----"
## [1] "Y[ 7000 ]: 1 0 0 0 1 0 0 1 ( 137 )"
## [1] "predict Y~: 1 0 0 0 1 0 0 1 ( 137 )"
## [1] "-----"
```

In fact, from the 6000th summation on, all the printed estimates are in fact correct.

## 4 Results

The eventual purpose is to use the model generated by the `trainr()` function as an input to the `predictr()` function, in order to predict the values of new inputs.

Table 11: Test Data

```
# create test inputs
C1 = int2bin( sample(0:127, 7000, replace=TRUE) )
C2 = int2bin( sample(0:127, 7000, replace=TRUE) )
```

Now predict using the `predictr()` function.

Table 12: `predictr()`

```
# predict
D <- predictr(m1,
              C1,
              C2,
              binary_dim = 8,
              alpha      = 0.1,
              input_dim  = 2,
              hidden_dim = 10,
              output_dim = 1 )
```


We can now convert the predictions and the inputs back to decimals and plot them.

Figure 2: Evaluate Predictions

```
# convert back to decimal
C1 <- bin2int(C1)
C2 <- bin2int(C2)
D <- bin2int(D)

# inspect the differences
qplot( D-(C1+C2) )

## 'stat_bin()' using 'bins = 30'. Pick better value with
'binwidth'.
```



As can be seen from the results, the difference is almost always 0.

## 5 Conclusion

CRAN and the rest of the R ecosystem show that there is a strong interest in using the R language for neural network analysis. Existing package such as the built in nnet package and the caret package make available very powerful neural network tools to R users. The RSNNs package acts as an R wrapper for the Stuttgart Neural Network Simulator library, which is written in C, and thereby makes available to partial RNNs such as Elman and Jordan networks.

The enormous popularity of full Recurrent Neural Networks in other languages, primarily Python and C, show that there is a great amount of interest for using this methodology, including interest from Economist, Data Scientists, and other non-professional programmers. Although Python is a relatively accessible programming language for laymen, it has a smaller user base in terms

of data analysts. The `rnn` package attempts to address this need by showing that Recurrent Neural Networks can be made available and perhaps more importantly, made available in native R, which allows user to delve into the code and understand the method and developer a more thorough understanding of how to use it.

## A Source code of trainr() function

0\_Users\_quast\_rnn\_R\_trainr.R

```

1 #' @name trainr
2 #' @export
3 #' @importFrom stats runif
4 #' @title Recurrent Neural Network
5 #' @description Trains a Recurrent Neural Network.
6 #' @param Y vector of output values
7 #' @param X1 vector of input values
8 #' @param X2 vector of input values
9 #' @param binary_dim dimension of binary representation
10 #' @param alpha size of alpha
11 #' @param input_dim dimension of input layer, i.e. how many numbers
    to sum
12 #' @param hidden_dim dimension of hidden layer
13 #' @param output_dim dimension of output layer
14 #' @param print should train progress be printed
15 #' @return a model to be used by the predictr function
16 #' @examples
17 #' # create training numbers
18 #' X1 = sample(0:127, 7000, replace=TRUE)
19 #' X2 = sample(0:127, 7000, replace=TRUE)
20 #'
21 #' # create training response numbers
22 #' Y <- X1 + X2
23 #'
24 #' # convert to binary
25 #' X1 <- int2bin(X1)
26 #' X2 <- int2bin(X2)
27 #' Y <- int2bin(Y)
28 #'
29 #' # train the model
30 #' trainr(Y,
31 #'       X1,
32 #'       X2,
33 #'       binary_dim = 8,
34 #'       alpha      = 0.1,
35 #'       input_dim  = 2,
36 #'       hidden_dim = 10,
37 #'       output_dim = 1,
38 #'       print = 'full ' )
39 #'
40
41
42 trainr <- function(Y, X1, X2, binary_dim, alpha, input_dim, hidden_
    dim, output_dim, print = c('none', 'minimal', 'full')) {
43
44   # check what largest possible number is
45   largest_number = 2^binary_dim
46
47   # initialize neural network weights
48   synapse_0 = matrix(stats::runif(n = input_dim*hidden_dim, min=-1,
    max=1), nrow=input_dim)
49   synapse_1 = matrix(stats::runif(n = hidden_dim*output_dim, min
    =-1, max=1), nrow=hidden_dim)

```

```

50 synapse_h = matrix(stats::runif(n = hidden_dim*hidden_dim, min
    =-1, max=1), nrow=hidden_dim)
51
52 synapse_0_update = matrix(0, nrow = input_dim, ncol = hidden_dim)
53 synapse_1_update = matrix(0, nrow = hidden_dim, ncol = output_dim
    )
54 synapse_h_update = matrix(0, nrow = hidden_dim, ncol = hidden_dim
    )
55
56 # training logic
57 for (j in 1:dim(Y)[1]) {
58
59     if(print != 'none' && j %% 1000 == 0) {
60         print(paste('Summation_number:', j))
61     }
62
63     # generate a simple addition problem (a + b = c)
64     a = X1[j,]
65     b = X2[j,]
66
67     # true answer
68     c = Y[j,]
69
70     # where we'll store our best guess (binary encoded)
71     d = matrix(0, nrow = 1, ncol = binary_dim)
72
73     overallError = 0
74
75     layer_2_deltas = matrix(0)
76     layer_1_values = matrix(0, nrow=1, ncol = hidden_dim)
77     # layer_1_values = rbind(layer_1_values, matrix(0, nrow=1, ncol
    =hidden_dim))
78
79     # moving along the positions in the binary encoding
80     for (position in 0:(binary_dim-1)) {
81
82         # generate input and output
83         X = cbind(a[binary_dim - position], b[binary_dim - position])
84         y = c[binary_dim - position]
85
86         # hidden layer (input + prev hidden)
87         layer_1 = sigmoid((X%%synapse_0) + (layer_1_values[dim(layer
    _1_values)[1],] %% synapse_h))
88
89         # output layer (new binary representation)
90         layer_2 = sigmoid(layer_1 %% synapse_1)
91
92         # did we miss?... if so, by how much?
93         layer_2_error = y - layer_2
94         layer_2_deltas = rbind(layer_2_deltas, layer_2_error *
    sigmoid_output_to_derivative(layer_2))
95         overallError = overallError + abs(layer_2_error)
96
97         # decode estimate so we can print it out
98         d[binary_dim - position] = round(layer_2)
99
100        # store hidden layer so we can print it out

```

```

101     layer_1_values = rbind(layer_1_values, layer_1)
102
103     if(print == 'full' && j %% 1000 == 0) {
104         print(paste('x1:', a[binary_dim - position]))
105         print(paste('x2:', b[binary_dim - position], '+'))
106         print('_____')
107         print(paste('y:~', c[binary_dim - position]))
108         print(paste('y^:', d[binary_dim - position]))
109         print('=====')
110     }
111 }
112
113 future_layer_1_delta = matrix(0, nrow = 1, ncol = hidden_dim)
114
115 for (position in 0:(binary_dim-1)) {
116
117     X = cbind(a[position+1], b[position+1])
118     layer_1 = layer_1_values[dim(layer_1_values)[1]-position,]
119     prev_layer_1 = layer_1_values[dim(layer_1_values)[1]-(
120         position+1),]
121
122     # error at output layer
123     layer_2_delta = layer_2_deltas[dim(layer_2_deltas)[1]-
124         position,]
125     # error at hidden layer
126     layer_1_delta = (future_layer_1_delta %*% t(synapse_h) +
127         layer_2_delta %*% t(synapse_1)) *
128         sigmoid_output_to_derivative(layer_1)
129
130     # let's update all our weights so we can try again
131     synapse_1_update = synapse_1_update + matrix(layer_1) %*%
132         layer_2_delta
133     synapse_h_update = synapse_h_update + matrix(prev_layer_1) %*
134         % layer_1_delta
135     synapse_0_update = synapse_0_update + t(X) %*% layer_1_delta
136
137     future_layer_1_delta = layer_1_delta
138 }
139
140 synapse_0 = synapse_0 + ( synapse_0_update * alpha )
141 synapse_1 = synapse_1 + ( synapse_1_update * alpha )
142 synapse_h = synapse_h + ( synapse_h_update * alpha )
143
144 synapse_0_update = synapse_0_update * 0
145 synapse_1_update = synapse_1_update * 0
146 synapse_h_update = synapse_h_update * 0
147
148 # convert d to decimal
149 out = b2i(as.vector(d))
150
151 # print out progress
152 if(print != 'none' && j %% 1000 == 0) {
153     print(paste('Error:', overallError))
154     print(paste('X1[', j, ']:', paste(a, collapse = '~'), '~', '(
155         ', b2i(a), ')'))
156     print(paste('X2[', j, ']:', paste(b, collapse = '~'), '+', '(
157         ', b2i(b), ')'))

```

```

151     print('=====')
152     print(paste('Y[', j, ']:', paste(c, collapse = '_'), '_', '(
      ', b2i(c), ')'))
153     print(paste('predict_Y^:', paste(d, collapse = '_'), '_', '
      (' , out, ')'))
154     print('=====')
155   }
156 }
157
158 # output object with synapses
159 return(list(synapse_0 = synapse_0, synapse_1 = synapse_1, synapse
      _h = synapse_h))
160
161 }

```

## B Source code of predictr() function

```

1_Users_quast_rnn_R_predictr.R
1 #' @name predictr
2 #' @export
3 #' @importFrom stats runif
4 #' @title Recurrent Neural Network
5 #' @description Trains a Recurrent Neural Network.
6 #' @param model output of the trainr function
7 #' @param X1 vector of input values
8 #' @param X2 vector of input values
9 #' @param binary_dim dimension of binary representation
10 #' @param alpha size of alpha
11 #' @param input_dim dimension of input layer, i.e. how many numbers
      to sum
12 #' @param hidden_dim dimension of hidden layer
13 #' @param output_dim dimension of output layer
14 #' @param print should train progress be printed
15 #' @return vector of predicted values
16 #' @examples
17 #' # create training numbers
18 #' X1 = sample(0:127, 7000, replace=TRUE)
19 #' X2 = sample(0:127, 7000, replace=TRUE)
20 #'
21 #' # create training response numbers
22 #' Y <- X1 + X2
23 #'
24 #' # convert to binary
25 #' X1 <- int2bin(X1)
26 #' X2 <- int2bin(X2)
27 #' Y <- int2bin(Y)
28 #'
29 #' # train the model
30 #' m1 <- trainr(Y,
31 #'             X1,
32 #'             X2,
33 #'             binary_dim = 8,
34 #'             alpha = 0.1,
35 #'             input_dim = 2,

```



```

36 #'          hidden_dim = 10,
37 #'          output_dim = 1  )
38 #'
39 #' # create test inputs
40 #' A1 = int2bin( sample(0:127, 7000, replace=TRUE) )
41 #' A2 = int2bin( sample(0:127, 7000, replace=TRUE) )
42 #'
43 #' # predict
44 #' B <- predictr(m1,
45 #'               A1,
46 #'               A2,
47 #'               binary_dim = 8,
48 #'               alpha      = 0.1,
49 #'               input_dim  = 2,
50 #'               hidden_dim = 10,
51 #'               output_dim = 1  )
52 #'
53 #' # convert back to integers
54 #' A1 <- bin2int(A1)
55 #' A2 <- bin2int(A2)
56 #' B <- bin2int(B)
57 #'
58 #' # inspect the differences
59 #' table( B-(A1+A2) )
60 #'
61 #' # plot the difference
62 #' hist( B-(A1+A2) )
63 #'
64
65
66 predictr <- function(model, X1, X2, binary_dim, alpha, input_dim,
67                       hidden_dim, output_dim, print = c('none', 'minimal', 'full')) {
68
69   # check what largest possible number is
70   largest_number = 2^binary_dim
71
72   # create output vector
73   Y <- matrix(nrow = dim(X1)[1], ncol = binary_dim)
74
75   # load neural network weights
76   synapse_0 = model$synapse_0
77   synapse_1 = model$synapse_1
78   synapse_h = model$synapse_h
79
80   # training logic
81   for (j in 1:dim(X1)[1]) {
82     if(print != 'none' && j %% 1000 == 0) {
83       print(paste('Summation_number:', j))
84     }
85
86     # generate a simple addition problem (a + b = c)
87     a = X1[j,]
88     b = X2[j,]
89
90     # where we'll store our best guess (binary encoded)
91     d = matrix(0, nrow = 1, ncol = binary_dim)

```

```

92
93 overallError = 0
94
95 layer_2_deltas = matrix(0)
96 layer_1_values = matrix(0, nrow=1, ncol = hidden_dim)
97 # layer_1_values = rbind(layer_1_values, matrix(0, nrow=1, ncol
    =hidden_dim))
98
99 # moving along the positions in the binary encoding
100 for (position in 0:(binary_dim-1)) {
101
102     # generate input and output
103     X = cbind(a[binary_dim - position], b[binary_dim - position])
104
105     # hidden layer (input ~+ prev_hidden)
106     layer_1 = sigmoid((X%%synapse_0) + (layer_1_values[dim(layer
        _1_values)[1],] %% synapse_h))
107
108     # output layer (new binary representation)
109     layer_2 = sigmoid(layer_1 %% synapse_1)
110
111     # decode estimate so we can print it out
112     d[binary_dim - position] = round(layer_2)
113
114     # store hidden layer so we can print it out
115     layer_1_values = rbind(layer_1_values, layer_1)
116
117     if(print == 'full' && j %% 1000 == 0) {
118         print(paste('x1:', a[binary_dim - position]))
119         print(paste('x2:', b[binary_dim - position], '+'))
120         print('_____')
121         print(paste('y^:', d[binary_dim - position]))
122         print('=====')
123     }
124 }
125
126 # output to decimal
127 out = b2i(as.vector(d))
128
129 # print out progress
130 if(print != 'none' && j %% 1000 == 0) {
131     print(paste('Error:', overallError))
132     print(paste('X1[', j, ']:', paste(a, collapse = '\n'), '\n', '(
        ', b2i(a), ')'))
133     print(paste('X2[', j, ']:', paste(b, collapse = '\n'), '\n', '(
        ', b2i(b), ')'))
134     print('_____')
135     print(paste('predict_Y^:', paste(d, collapse = '\n'), '\n', '
        (' , out, ')'))
136     print('=====')
137 }
138
139 # store value
140 Y[j,] <- d
141 }
142
143 # return output vector

```

```
144 |   return( Y )  
145 | }
```